

The Statistics of Causal Inference: A View from Political Methodology

Luke Keele

Department of Political Science, 211 Pond Lab, Penn State University, University Park, PA 19130
e-mail: ljk20@psu.edu (corresponding author)

Edited by R. Michael Alvarez

Many areas of political science focus on causal questions. Evidence from statistical analyses is often used to make the case for causal relationships. While statistical analyses can help establish causal relationships, it can also provide strong evidence of causality where none exists. In this essay, I provide an overview of the statistics of causal inference. Instead of focusing on specific statistical methods, such as matching, I focus more on the assumptions needed to give statistical estimates a causal interpretation. Such assumptions are often referred to as identification assumptions, and these assumptions are critical to any statistical analysis about causal effects. I outline a wide range of identification assumptions and highlight the design-based approach to causal inference. I conclude with an overview of statistical methods that are frequently used for causal inference.

1 Introduction

One central task of the scientific enterprise is establishing causal relationships. Take one example from the comparative politics literature. One well-known finding is that democracies are less likely to engage in the repression of human rights (Poe and Tate 1994). We can just treat this as a descriptive finding: democratic governance is correlated with lower levels of repression. This descriptive finding, however, begs a causal question: if a country becomes more democratic, will it then engage in less repression? Rarely are we content with statistical associations. Instead, we often seek to establish causal relationships.

Causality is something we all understand, since we use it in our daily life. It refers to the relational concept where one set of events causes another. Causal inference is the process by which we make claims about causal relationships. While causality seems a simple concept in everyday life, the establishment of causal relationships in many contexts is a difficult enterprise. Early models of causality focused on unique causes such as gravity. Gravity always causes things to fall to the earth and is the unique cause of that action. In biological and social applications, outcomes rarely have unique causes, as causes tend to be contingent. In such contexts, the counterfactual model of causality is useful. Under the counterfactual model, rather than define causality purely in terms of observable events, causation is defined in terms of observable and unobservable events. Thus, I say, if Iraq had been democratic, war would not have broken out. This is a counterfactual statement about the world that asserts that if a cause had occurred an effect would have followed. This counterfactual approach is based on the idea that some of the information needed to make a causal inference is unobserved and thus some assumptions must be made before I can make a causal inference.¹

Authors' note: For comments I thank the editors and the four anonymous reviewers. I also thank Rocío Titiunik, Jasjeet Sekhon, Paul Rosenbaum, and Dylan Small for many insightful conversations about these topics over the years. In the online Supplementary Materials, I provide further information about software tools to implement many of the methodologies discussed in this essay. Supplementary materials for this article are available on the *Political Analysis* Web site.

¹See Hidalgo and Sekhon (2011) for an overview of different models of causality and the rise of the counterfactual model.

In the social sciences, data and statistical analyses are often used to test causal claims. Over the last 20 years, the potential outcomes framework, a manifestation of the counterfactual model of causality, has come to dominate statistical thinking about causality. What is behind the popularity of this approach? Why do some, myself included, view this framework as an improvement over the past and not simply a “re-labeling” of existing statistical concepts? First, the counterfactual approach has provided new insights into the assumptions needed for data to be informative about causality. Specifically, there has been a renewed interest in the assumptions needed for causal inference and unpacking the exact meaning of those assumptions. Second, there has been a renewed emphasis on research design and the design-based approach. As I discuss below, the phrase “design-based approach” does not have universal definition, but there is widespread agreement that statistical analyses are more convincing when the research design has been carefully constructed to bolster assumptions before estimation.

In this essay, I provide a roadmap to the statistics of causal inference. I divide the statistics of causal inference into three parts: causal identification, the design-based approach, and statistical tools. I begin with an introduction to the concept of causal identification and identification analyses. An identification analysis identifies the assumptions needed for statistical estimates to be given a causal interpretation. Next, a researcher must select an identification strategy or research design. In this section, I also provide a brief overview of several common identification strategies.

Once an identification strategy has been selected, the analyst can often use elements of the design-based approach to improve the research design. The design-based approach is a set of techniques that can make identification more credible without the use of parametric statistical models and without using outcomes. Finally, I provide a brief overview of statistical tools like matching and inverse probability (IP) weighting, which are commonly used for the final part of a causal analysis: estimation of treatment effects. In this section, I also review how the mode of statistical inference changes when the focus is on causal effects. Through this structure, I clarify the distinctions among identification, design, and statistical analysis.

2 Identification

I begin with a summary of identification, which is an extremely important concept in causal inference. One way to describe whether a statistical estimate can be given a causal interpretation is to discuss whether its target causal estimand, defined below, is identified or not. Identification concepts are invoked (often implicitly) in any analysis that purports to present a causal effect. Confusion often develops, however, since the concept of identification is more general. Identification problems arise in a number of different settings in statistics. I start with a general discussion of identification, but then outline the specific identification problem that underlies causal inference.

2.1 Basics of Identification

Informally, we say a parameter in a model is identified if it is theoretically possible to learn the true value of that parameter with an infinite number of observations (Matzkin 2007, sec. 3.1). Conversely, for problems of identifiability, there are cases where even if we have an infinite number of observations, we don’t have enough information to learn about the true value of a parameter in the model. Manski (1995) separates the problem of inference into two components: identification and statistical. Under the identification part of inference, we seek to describe the conclusions that can be drawn with an infinite sample. If identification fails, nothing can be learned even if the sample is infinite. Studies of statistical inference, on the other hand, focus on what can be learned with finite samples.

There are many identification problems in statistics. For example, studies of ecological inference are based on an identification problem, where one attempts to identify the parameters of mixtures of probability distributions using only knowledge of the marginal distributions. This type of identification problem occurs when we attempt to make inferences about units based on aggregates such

as inferences about voters based on aggregating voting data. Inferences based on missing data form a different identification problem. Causal inference is yet another identification problem. Importantly, the causal inference identification problem can only be resolved through assumptions, which is not always the case for other identification problems. Consider the identification problem created by missing data. We can solve that identification problem using a set of assumptions about the missing data. Alternatively, we might alter the data collection process such that no data are missing, thus avoiding the use of assumptions. Under causal inference, there is no alternative method to resolve the identification problem other than through assumptions, since certain counterfactual quantities are unobservable.

To understand whether causal identification holds, we must perform an *identification analysis*. In an identification analysis, we consider whether it is possible to compute a causal effect from data with an infinite sample. In an identification analysis, the analyst provides both a formal statement of the set of assumptions needed to identify a particular causal effect and a proof that those assumptions lead to an identified causal effect. For example, one could state and prove the assumptions that must hold to compute a causal effect from a randomized experiment conducted with an infinite sample. More commonly one invokes an existing identification analysis and stipulates identification under that set of assumptions.

Often analysts perform a *nonparametric* identification analysis. In a nonparametric identification analysis, we formally prove which non-model-based (functional form) assumptions are needed for identification.² Nonparametric identification analyses are considered important since they allow one to state the weakest set of assumptions needed for identification. Nonparametric identification assumptions are the assumptions that must hold to compute a causal effect with some hypothetical set of data that is infinite in size and without reference to any specific statistical model. As I will outline later, nonparametric identification often leads to a preference for nonparametric estimation methods.

Next, I consider the source of the identification problem in causal inference. The potential outcomes framework (see, e.g., Rubin 1974), often referred to as the Rubin Causal Model (RCM) (Holland 1986), is one way to formalize the causal inference identification problem. The RCM is the dominant model of causality in statistics at the moment. Like all models it is wrong, but it is also quite useful. In fact, the RCM is not the only model of causality that is embedded within a statistical framework. Dawid (2000) develops a decision theoretic approach to causality that rejects counterfactuals. Pearl (1995, 2009a) advocates for a model of causality based on nonparametric structural equations and path diagrams.

In the potential outcomes model, each unit has multiple potential outcomes but only one actual outcome. Potential outcomes represent unit-level behavior in the presence or absence of an intervention or treatment, and the actual outcome depends on actual treatment received. I denote a binary treatment with $D_i \in \{0, 1\}$, though the treatment need not be binary. The potential outcomes are Y_{iD} . The actual outcome is a function of treatment assignment and potential outcomes such that $Y_i = D_i Y_{i1} + (1 - D_i) Y_{i0}$. Under this framework, we can define various forms of the unit-level causal effect of D_i , which are comparisons of unit-level potential outcomes. One possible comparison is a difference in potential outcomes, $Y_{i1} - Y_{i0}$, but in general the comparisons can take different forms, such as a ratio: Y_{i1}/Y_{i0} .

We cannot estimate this unit-level causal effect since we do not observe the potential outcomes. The potential outcomes model formalizes the idea that the individual-level causal effect of D_i is unobservable, which is sometimes called the *fundamental problem of causal inference* and encapsulates the identification problem we face as causal inferences are based on comparisons of counterfactual quantities that can't be observed (Holland 1986). In general, we focus on the average treatment effect (ATE):

$$\text{ATE} = \mathbb{E}[Y_{i1} - Y_{i0}]. \quad (1)$$

²Linearity and additivity, for example, are model-based functional form assumptions.

This is known as a causal estimand, since it based on a contrast of potential outcomes. It is separate from a statistical estimator or a specific point estimate that would be derived from observable data. In an identification analysis, we seek to identify specific estimands. The ATE is the average difference in the pair of potential outcomes averaged over the entire population of interest. Often causal estimands are defined as averages over specific subpopulations. For example, we might average over subpopulations defined by pretreatment covariates such as sex and estimate the ATE for females only. When the estimand is defined for a specific subpopulation, it is said to be more local. Frequently, the ATE is defined for the subpopulation exposed to the treatment or the ATE on the treated (ATT):

$$ATT = E[Y_{i1} - Y_{i0} | D_i = 1]. \quad (2)$$

Finally, we can define the relevant subpopulation in terms of potential outcomes. As I discuss later, the most well-known estimand to be defined in terms of a subpopulation based on potential outcomes comes from instrumental variables (IVs).³

I should note that I have followed common practice and written the estimands as averages. Causal identification rarely implies that only the middle (as represented by an average) of the treated and control distributions will differ. Analysts should always consider that causal effects might only be apparent at particular quantiles. I revisit this topic later when I consider methods of inference for causal effects.

For all these estimands, we face an identification problem, since there are terms in the estimand that are unobservable. Even if we had samples of infinite size, we still could not estimate the average causal effect without observing both potential outcomes. Using potential outcomes, we can clearly elucidate the unobservable quantities in the ATE estimand. I define π as the proportion of the sample assigned to the treatment condition. Using π , I can decompose the true ATE as a function of potential outcomes as follows:

$$\begin{aligned} E[Y_{i1} - Y_{i0}] \\ = \pi\{E[Y_{i1} | D_i = 1] - E[Y_{i0} | D_i = 1]\} + (1 - \pi)\{E[Y_{i1} | D_i = 0] - E[Y_{i0} | D_i = 0]\}. \end{aligned} \quad (3)$$

In equation (3), the ATE is a function of five quantities. Without additional assumptions we can estimate only three of those quantities directly from observed data. We can estimate π using $E[D_i]$. We can also readily estimate $E[Y_{i1} | D_i = 1]$ and $E[Y_{i0} | D_i = 0]$ using $E[Y_i | D_i = 1]$ and $E[Y_i | D_i = 0]$. However, we cannot estimate $E[Y_{i1} | D_i = 0]$ and $E[Y_{i0} | D_i = 1]$ from the data without assumptions. One is the average outcome under treatment for those units in the control condition, and the other is the average outcome under control for those in the treatment condition. That is, we face an identification problem, since these two quantities are unobserved counterfactuals, and no additional amount of data will allow us to estimate these quantities. Therefore, we must find a set of assumptions that allow for identification.

In causal inference, identification generally rests on the assumption that treatment status is independent of potential outcomes. Formally, this assumption is

$$Y_{i1}, Y_{i0} \perp D_i. \quad (4)$$

Why does this assumption identify causal effects? The expectation of the observed outcome conditional on $D_i = 1$ can be written as

$$\begin{aligned} E[Y_i | D_i = 1] &= E[Y_{i0} + D_i(Y_{i1} - Y_{i0}) | D_i = 1] \\ &= E[Y_{i1} | D_i = 1] \\ &= E[Y_{i1}], \end{aligned} \quad (5)$$

³Here, I implicitly invoke the SUTVA, which permits the assumption that we are actually observing the potential outcomes associated with each treatment condition. I discuss SUTVA in more detail later in this section.

where the last step follows from the assumption of independence. That is, taking the expectation of the observed outcome provides the expectation of the *potential* outcome when independence holds. Independence between treatment status and potential outcomes allows us to connect the observed outcomes to the potential outcomes. By the same logic, $\mathbb{E}[Y_i|D_i = 0] = \mathbb{E}[Y_{i0}]$. It follows from the above statements that

$$\begin{aligned} \text{ATE} &= \mathbb{E}[Y_{i1} - Y_{i0}] \\ &= \mathbb{E}[Y_{i1}] - \mathbb{E}[Y_{i0}] \\ &= \mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0]. \end{aligned} \tag{6}$$

That is, under the assumption of independence, the expectation of the unobserved potential outcomes is equal to the conditional expectations of the observed outcomes conditional on treatment assignment. The independence assumption allows us to connect unobservable potential outcomes to observed quantities in the data, though one additional assumption is also needed. I outline that additional assumption in the next section. When are we justified in assuming that independence holds between the treatment and the potential outcomes? I take up that question in the next section on identification strategies.

While independence between treatment and potential outcomes is one assumption often used for identification, typically additional assumptions are necessary for identification. Typically, we must also assume that the stable unit treatment value assumption (SUTVA) holds (Rubin 1986). SUTVA is made up of the two following components: (1) there are no hidden forms of treatment, which implies that for unit i under $D_i = d$, we assume that $Y_{id} = Y_i$ and (2) a subject's potential outcome is not affected by other subjects' exposure to the treatment.

The first component of SUTVA is often referred to as the consistency assumption in the epidemiological literature, and under this assumption we assume that for units exposed to a treatment we observe the potential outcomes for that treatment. The consistency assumption is somewhat controversial. Hernán and VanderWeele (2011) argue that the consistency assumption must be evaluated by analysts since it links observed data to the counterfactual outcomes. They argue that in the absence of consistency, one would not know which counterfactual contrast is being estimated by the data, which makes it difficult to base decision-making on a causal analysis. For example, if the treatment were "15 min of exercise," there are many different forms of exercise. They contend that it will be difficult to justify any decision-making based on effect estimates since we may not know which form of exercise actually made the treatment effective. In contrast, van der Laan, Haight, and Tager (2005) say that consistency is an axiom which can be taken for granted, while Pearl (2010) maintains that consistency immediately follows so long as the causal model is correct. In some sense, there are elements of truth to both sides. If potential outcomes are independent of the exercise treatment, we can rule out the presence of other causes. However, generating policy recommendations about this treatment may be difficult given the fact that the treatment may contain a large number of components.

The second part of the SUTVA assumption tends to be a more serious problem in many social science settings. The problem is that if we treat a unit and that unit can then spread some of that treatment to a control unit or units, the comparison is no longer between treated and control, but between a treated unit and partially treated unit. If one specifies a model of contagion for how the treatment spreads, one can make some progress toward identification, but if we have no knowledge of treatment spillovers, causal parameters will not be identified. Taking interference into account is currently a very active area of research in both the social sciences and statistics (Sinclair, McConnell, and Green, 2012; Tchetgen and VanderWeele 2012; Bowers, Fredrickson, and Panagopoulos 2013). Treatments that vary over time can also lead to SUTVA violations, as well as other complications. There has been considerable focus on treatments that vary over time in the biostatistics literature, and I suspect time-varying treatments will eventually be a topic of interest in political science (Robins 1997, 1999).

Next, I briefly review the two most common threats to identification. The first is confounding due to a common cause. Figure 1a contains a causal diagram of confounding. Confounding in




Fig. 1 Two threats to identification: confounding and selection.

many statistical texts is referred to as a spurious relationship. In this diagram, we might think that D_i is a cause of Y_i , but in fact L_i is a cause of both, while D_i is actually independent of Y_i . However, if we estimate the statistical association between D_i and Y_i , we will find them spuriously correlated, as represented by the dashed line. Concerns about confounding then are concerns about identifiability.

The next threat to identification stems from when D_i and Y_i both condition on a common effect. Rosenbaum (1984) identified this threat to identification as conditioning on a post-treatment covariate. Figure 1b represents this situation. The box around L_i represents conditioning. Often this is defined as selection since it arises from selection on the dependent variable. Selection tends to be a more subtle problem than confounding, since it can cause a failure of identification even when the treatment is independent of potential outcomes and confounding is ruled out. See Elwert and Winship (2014) for an excellent discussion of how threats to identification of this form can arise.

In sum, causal inference is based on an unavoidable identification problem. The first step in a causal analysis is the adoption of an identification strategy, a strategy for solving the causal inference identification problem. I turn to identification strategies next.

3 Identification Strategies

An identification strategy is simply a research design intended to solve the causal inference identification problem (Angrist and Pischke 2010).⁴ Part of an identification strategy is an assumption or set of assumptions that will identify the causal effect of interest. To ask what is your identification strategy is to ask what research design (and assumptions) one intends to use for the identification of a causal effect. The following review of identification strategies is necessarily brief, but I highlight those most commonly used. Readers interested in a more in-depth review from different perspectives should consult Angrist and Pischke (2009), Morgan and Winship (2014), and Rosenbaum (2010).

3.1 Randomized Experiments

The randomized experiment is often considered the “gold standard” among identification strategies. Here, subjects are assigned to D_i via some random mechanism like the toss of a fair coin. The typical estimand in a basic randomized design is the ATE, which under this identification strategy is equivalent to the ATT. Of course, randomization does not necessarily imply that only averages will differ across the treated and control groups. As I discuss later, other features of the treated and control distributions may also be of interest.

⁴I should note that there are various inconsistencies in how the phrase identification strategy is defined even within the same set of authors since Angrist and Pischke (2009) define an identification strategy as “the manner in which a researcher uses observational data to approximate a real experiment.” This definition appears to only include non-experimental research designs as identification strategies, which seems overly narrow.

See Rosenbaum (2010, chap. 2) and Gerber and Green (2012) for details on experiments. Here, I want to convey what is special about this identification strategy. The key strength of experiments is that the researcher has the ability to impose independence between treatment status and potential outcomes on a set of units because he or she can impose a particular type of assignment process. As I outlined above, if the treatment is independent of the potential outcomes, then the treatment effect parameter is identified. Short of incorrectly generating random treatment assignments, under this identification strategy the analyst knows that independence holds, which allows the researcher to assert that the treated and control groups will be identical in all respects, observable and *unobservable*, save receipt of the treatment with arbitrarily high probability as the sample size grows large. This implies that randomization allows us to rule out confounding due to a common cause.

Of course, randomization is not a cure-all. We must assume SUTVA holds. Moreover, experiment may not give valid causal estimates when attrition is present. Attrition is when subject outcomes are not available after randomization, and this missingness is correlated with treatment status. Another complication in experiments is noncompliance. It is often the case that subjects do not comply with their assigned treatment status. A full discussion of noncompliance and attrition is beyond the scope of this article. See Gerber and Green (2012) for a detailed discussion of both topics. Later I discuss noncompliance in more detail, since it forms a separate identification strategy.

Finally, a randomized experiment identifies the treatment effect within the population used in the study. This treatment effect may or may not extrapolate to other populations. To ensure valid extrapolation, one either needs random sampling in addition to randomization of treatment or additional assumptions. Given this fact, experiments are often said to be internally valid, but they may lack external validity (Campbell and Stanley 1963). Whether this is a feature or a bug is a matter of substantial disagreement. Many of those who label themselves as interested in causal inference tend to value internal validity over external validity. If our concern is observing a causal effect, we might place more value on a well-executed laboratory experiment than an observed association from a very large representative sample of data. As we will see, some identification strategies explicitly work based on comparisons of comparable but unrepresentative subpopulations. I explain the logic behind the value placed on internal validity in the next section.

3.2 *Natural Experiments*

The next identification strategy is based on natural experiments. A natural experiment is a real-world situation that produces haphazard assignment to a treatment (Rosenbaum 2010, 67). The hope is that a natural intervention will create as-if randomized treatment assignment and thereby produce independence between treatment assignment and unit level potential outcomes. Of course, randomization in an experiment is a fact, while haphazard treatment assignment often requires considerable judgment to justify it as as-if random. The circumstances of the natural experiment speak to whether the claim of as-if random assignment is credible, but there is no way to know whether assignment is as good as randomized. An example is helpful.

Lyall (2009) seeks to understand whether indiscriminate violence increases insurgent attacks. To that end, he exploits shelling patterns by Russian troops in Chechnya that appear to be at worst indiscriminate and at best as-if random. He does find that the treatment, being shelled, appears to be uncorrelated with pre-treatment covariates, as would be the case in a randomized experiment. The difficulty is that, unlike with randomization, we don't know whether the patterns are truly random since they are beyond the control of the analyst. As such, natural experiments often require careful justification for the as-if random nature of assignment. The basic template, however, is present in the study by Lyall (2009). He finds a real-world situation that appears to mimic a randomized experiment. Exploiting such circumstances is often a very credible identification strategy. Like randomized experiments, the focus is on internal validity. We have no way of knowing whether the causal effect in Chechnya would hold in another circumstance, but what we hope to observe is a causal effect operating in relative isolation from the very real threats of confounding.

3.3 Instrumental Variables

Informally, an instrument is a random push to accept a treatment, but the push can only affect the outcome if it induces units to take the treatment. Holland (1988) outlined the randomized encouragement design as the prototype of an instrument. He described this design as an experiment where some participants are encouraged to exercise. While subjects are randomly encouraged to exercise, subjects then select their exposure to the exercise treatment in that they select whether to exercise or not. Moreover, some of those assigned to the non-exercise arm will decide to exercise. Later all participants are measured on the outcome.

There are two effects of interest in designs of this type. In this design, the effect of being assigned to encouragement is identified since this has been randomly assigned. This estimand is often called the intention-to-treat (ITT) effect. This estimand tells us whether encouragement changes the outcome. Under additional assumptions, the method of IV identifies the effect of the treatment, exercise, as opposed to the effect of being assigned to exercise encouragement (Angrist, Imbens, and Rubin 1996). Specifically, IV identifies the average effect among those induced to take the treatment by a randomized encouragement. The IV estimand is often referred to as either the complier average causal effect (CACE) or the local ATE (LATE). The IV estimand is local since it is defined for a subpopulation: the compliers. However, this subpopulation is defined in terms of potential outcomes, since compliance status is unobservable for any particular unit (Angrist, Imbens, and Rubin 1996).

For IV to provide valid causal inferences, the five assumptions outlined by Angrist, Imbens, and Rubin (1996) must hold. The assumptions needed for the IV estimand to be identified are (1) ignorable (as-if random) assignment of the encouragement; (2) the SUTVA; (3) no direct effect of the instrument (here, encouragement) on the outcome also known as the exclusion restriction; (4) monotonicity; and (5) the instrument must have a nonzero effect on the treatment. The first two assumptions are identical to those needed to identify the ITT effect. The other three are additional assumptions needed to identify the CACE.

Real-life circumstances can create circumstances that mimic the randomized encouragement design. More broadly, we can define an instrument as a haphazard nudge to accept a treatment. Here, IV becomes identification strategy based on a type of natural experiment. Hansford and Gomez (2010) are one example of using IV as a natural experiment identification strategy. They seek to understand whether lower turnout reduces the vote share for the Democratic Party. They exploit the fact that rainfall appears to decrease turnout on election day; and use it as an as-if random discouragement for turnout. If rainfall is a valid instrument, this allows them to identify the local effect of turnout on vote share among the counties discouraged to vote by rain on election day. While the IV identification strategy can be credible, when used as a natural experiment it requires great care. See Bound, Jaeger, and Baker (1995) for one example of a fairly spectacular failure of IV. See Sovey and Green (2011) for a more detailed overview of the IV identification strategy.

One important insight that originated in the statistical literature on IVs was the role of implicit constant effect assumptions. Angrist, Imbens, and Rubin (1996) clearly demonstrated that regression-based IV estimates required an assumption that the effect of the treatment was constant across units. They showed that under the nonparametric potential outcomes framework such assumptions could be relaxed. This insight has led to closer examination of implicit constant effects assumptions in many other identification strategies.

3.4 Regression Discontinuity Designs

The regression discontinuity (RD) design is another identification strategy that is typically classified as a type of natural experiment. In an RD design, assignment of the binary treatment, D_i , is a function of a known continuous covariate, S_i , usually referred to as the *forcing variable* or the *score*. In the sharp RD design, treatment assignment is a deterministic function of the score, where all units with score less than a known cutoff, c , are assigned to the control condition ($D=0$) and all units above the cutoff are assigned to the treatment condition ($D_i=1$). In the fuzzy RD design, assignment to the treatment is a random variable given the score, but the probability of receiving

treatment conditional on the score, $P(D_i = 1|S_i)$, must jump discontinuously at c . This implies that it is possible for some units with scores below c to receive the treatment. The fuzzy RD design results in an equivalence between RD and IV (Hahn, Todd, and van der Klaauw 2001). See Lee and Lemieux (2010) for a much lengthier description of RD designs.

Hahn, Todd, and van der Klaauw (2001) demonstrate that for the sharp RD design to be identified the potential outcomes must be a *continuous* function of the score in the neighborhood around the discontinuity. Under this continuity assumption, the potential outcomes can be arbitrarily correlated with the score, so that, for example, people with higher scores might have higher potential gains from treatment. The continuity assumption is a formal statement of the idea that individuals very close to the cutoff but on opposite sides of it are comparable or good counterfactuals for each other. Thus, continuity of the conditional regression function is enough to identify the causal effect *at the cutoff*. The idea is that if nothing in the potential outcomes changes abruptly at the cutoff other than the probability of receiving treatment, any jump in the conditional expectation of the outcome variable at the cutoff is attributed to the effects of the treatment. Often it is assumed that there is a neighborhood around the cutoff where treatment status is considered as good as randomly assigned. Such an interpretation requires an additional identification assumption (Cattaneo, Frandsen, and Titiunik 2014). Here, the analyst must choose a neighborhood or window around the cutoff where treatment status is assumed to be as-if randomly assigned.

The RD design is another example of where the estimand changes as a function of the design. The RD design identifies a LATE for the subpopulation of individuals whose value of the score is at or near c , so the estimand is restricted to a subset of units on either side of the threshold that are thought to be good counterfactuals. In this design, it is only possible to identify the treatment effect among a small subpopulation around the cutoff. Here, complications can arise since unmodeled nonlinearity can be mistaken for a treatment effect. See Angrist and Pischke (2010) for an example of this.

Lee and Lemieux (2010) note that one strength of the RDD is that it is a design. Like randomization, some decision-maker must implement a treatment assignment mechanism based on a continuous score and a cutoff for a population of subjects. Lee and Lemieux (2010) emphasize this aspect of an RDD which distinguishes it from many natural experiments that rely on an instrument such as rainfall, which is certainly stochastic in some sense, but is not a controlled treatment assignment mechanism. Moreover, RD designs have gained further credibility by recovering experimental benchmarks (Cook, Shadish, and Wong 2008). However, see Caughey and Sekhon (2011) for one example where the assumptions of an RD design fail.

3.5 Selection on Observables

Under the “selection on observables” identification strategy, the analyst asserts that there is some set of covariates such that treatment assignment is random conditional on these covariates (Barnow, Cain, and Goldberger 1980). Under this assumption, there are no unobservable differences between the treated and control groups. This assumption has a number of different names, which include “conditional ignorability” and “no omitted variables.” All of these are statements of the same idea: we seek to make the treatment independent of the potential outcomes conditional on observed covariates. Under this identification strategy, we assume that the treatment is conditionally independent of potential outcomes. Critically, the selection on observables assumption is nonrefutable, insofar as it cannot be verified with observed data (Manski 2007).

Given this set of “correct” covariates, we can use statistical adjustment methods such as regression, matching, or weighting to make conditional independence hold. In regression terms, this implies that we tend to prefer longer specifications to shorter specifications. Of course, there are dangers in pursuing overly long specifications. While we need to include all covariates that predict the outcome and treatment, we cannot condition on any covariates that are affected by treatment (Rosenbaum 1984) without further assumptions. Even in a randomized experiment, conditioning on covariates that are affected by the treatment will bias our estimate of the treatment effect. This is sometimes known as over or bad control. See Angrist and Pischke (2009, 69) for an accessible review of the formal statement of the bias that arises from controlling for post-treatment covariates.

In an experiment, we can clearly delineate between the pre-treatment and post-treatment time periods. In observational data, that is often more difficult. In survey data, for example, it can be difficult to delineate any covariates as either pre- or post-treatment.

In a further complication, Pearl (2009a, 2009b) warns that adjustment for certain types of pre-treatment covariates can cause bias. This is known as “M-Bias” and arises from conditioning when there is a particular structure of unobservable covariates that create what is known as a “collider.” Ding and Miratrix (forthcoming) show that while M-Bias is generally small, there are rare cases where blind inclusion of pre-treatment covariates can induce severe bias. As such, one must choose specifications with some care. I can’t emphasize enough that selection on observables is a very strong assumption. It is often difficult to imagine that selection on observables is plausible in many contexts. Generally, selection on observables needs to be combined with a number of different design elements before it becomes plausible. I outline design elements in the next section.

3.6 *Selection on Observables with Temporal Data*

As I noted above, the selection on observables identification strategy requires that all differences between treated and control are observable. We can weaken this assumption when we observe units across multiple time periods. When there are data at multiple time periods, three different identification strategies are possible: fixed effects, differences-in-differences (DID), and identification based on lags. See Angrist and Pischke (2009, chap. 5) for a more in-depth overview of these related identification strategies.

Under the fixed effects identification strategy, if we use repeated observations on individuals, we assume that treatment is independent of potential outcomes so long as any confounders are time invariant. Therefore, if confounders are time invariant it doesn’t matter if they are unobserved. However, we must also assume that the treatment effect is linear and additive, which is a strong constraint on how units respond to treatment. DID is a second identification strategy based on repeated observations. Angrist and Pischke (2009, 228) describe DID as a fixed effects identification strategy using aggregate data. Here, the key identifying assumption is that trends in the outcome would be the same across treated and control groups in absence of the treatment. That is, we must assume that no other events beside the treatment alter the temporal path of either the treated or control groups.

The next identification strategy conditions on unobservables in an indirect fashion using past outcomes. Under this identification strategy, we assume selection on observables, but we condition on some number of lags of the outcome. Why is this an improvement over simply conditioning on observables? The key insight is that lagged outcomes are a function of both observable covariates and unobservables. As such, if we condition on lagged outcomes we can indirectly condition on unobservables. The method of synthetic case control relies on this identification strategy (Abadie and Gardeazabal 2003; Abadie, Diamond, and Hainmueller 2010).

While all these methods do allow for conditioning on unobservables, they all require the unobservables to have a very specific configuration. For all three strategies, the key assumption remains untestable. See Arceneaux, Gerber, and Green (2006) for one example of where identification based on lags fails. This should serve a useful reminder that identification under any version of the selection on observables assumption is fraught with uncertainty.

3.7 *Partial Identification*

The goal under most identification strategies is point identification—identification of a single parameter that describes the causal effect of D_i . An alternative approach is to instead place bounds on the treatment effect, which can typically be done with weaker assumptions. The method of partial identification is most closely linked to the work of Manski (1990, 1995). See Mebane and Poast (2013) and Keele and Minozzi (2012) for examples in political science. Under partial identification, the analyst acknowledges that there is a fundamental tension between the credibility of assumptions and the strength of conclusions. As such the analysis proceeds by starting with the no-assumption bounds and adding assumptions about the nature of treatment response or assignment. By adding

the assumptions individually, it allows one to observe exactly which assumption provides an informative inference. Assumptions can also be combined for sharper inferences.

The partial identification strategy can be very useful. The discipline of adding assumptions in a specific order and debating the credibility of those assumptions is an important exercise. Moreover, it can be applied to any identification strategy. Lee (2009) uses a partial identification approach for randomized experiments with missing outcome data. Balke and Pearl (1997) use partial identification to relax the monotonicity assumption and exclusion restriction under the IVs identification strategy. Finally, partial identification also underpins many forms of sensitivity analysis.

3.8 Mediation Analysis

The final identification strategy I outline is rather different from those above. In an analysis of causal effects, we can broadly define three types of effects: total, direct, and indirect effects. The total effect is equivalent to the ATE. In a mediation analysis, we seek to decompose the total effect into indirect and direct effects. One criticism of the total effect is that it cannot tell the analyst why the treatment works, only that it does or does not. In a mediation analysis, the analysts posit a causal mechanism which depends on M_i , known as a mediating variable, which occurs post-treatment and is assumed to be affected by the treatment. The causal mediation effect represents the indirect effect of the treatment on the outcome through the mediating variable (Pearl 2001; Robins 2003). While the indirect effect represents the effect of the treatment through M_i , the direct effect represents the effect of the treatment through all other possible mediators. The goal in a mediation analysis is to decompose the total effect into its indirect and direct components.

Identification in a mediation analysis proceeds in two parts. First, one makes the case for identifiability of the total effect. Identification of the direct and indirect effects requires an additional assumption. Typically analysts use an assumption known as sequential ignorability, which rules out confounding between M_i and Y_i (Imai et al. 2011). That is, the analysts must assume that all pre-treatment covariates that might confound the relationship between the mediator and outcome are observed. Thus, the focus here is on the identification assumptions for the indirect and direct effects, while identification of the total effect depends on one of the identification strategies listed above. As such, this identification strategy is generally secondary since one must first make a case for the identifiability of the total effect. If identifiability of the total effect is doubtful, there is little use in pursuing a mediation analysis.

3.9 Reasoning About Assumptions

Finally, I highlight one of the more important skills needed for causal inference. Critically, the plausibility of an identification strategy depends on the empirical context. For every identification strategy outlined above, one can find contexts where it is plausible and other contexts where that same strategy is indefensible.

Take the selection on observables identification strategy, which is generally viewed as the weakest identification strategy. Sekhon and Titiunik (2012) present an example of estimating incumbency effects based on the redistricting process, where selection on observables is credible. In their example, voters are assigned to an incumbency treatment in the redistricting process. They note that since we know that state legislators use observable data to decide how to draw districts, we have good reason to believe the treatment assignment process is observable. That is, selection to treatment is based on observables. Thus, redistricting makes selection on observables a plausible identification strategy. Take DID as a second example. Gordon (2011) is one example where a DID identification strategy is highly plausible. Alternatively, Keele and Minozzi (2012) outline an example where a DID identification strategy generally fails.

As such, reasoning about the plausibility of an identification strategy in a specific empirical context is a critical part of any statistical analysis that purports to be causal. Since untestable assumptions are unavoidable in causal inference, it is only through the careful understanding of those assumptions that one can make a case for their plausibility in a given context. As such, the researcher must think deeply about the assumptions and part of the analysis should be a

well-reasoned defense of the identification strategy. Qualitative information is often critical for defending the identification strategy. Reasoning about assumptions is often not part of a statistical analysis, but it must be when the goal is to identify causal effects.

A number of important contributions in the literature on causal inference stem from a re-articulation of identification assumptions in a way that allows for a better understanding of those assumptions. For example, Lee (2008) developed a useful way to interpret the continuity assumption in the RD design. He defines the score as $S_i = W_i + e_i$, where W_i represents efforts by agents to sort above and below c and e_i is a stochastic component. When e is small, this implies that agents are able to precisely sort around the threshold, treatment is mostly determined by self-selection, and identification is less plausible. However, when e_i is larger, agents will have difficulty self-selecting into treatment, and whether an agent is above or below the threshold is essentially random. This behavioral interpretation of the continuity assumption allows aids in the assessment of the RD design.

The rewriting of IVs using the potential outcomes framework is another example of how restating assumptions can be incredibly important. Angrist, Imbens, and Rubin (1996) took the traditional statement of IV assumptions based on covariance restrictions and restated them into a form that allows for better reasoning about their plausibility. Many of the mistakes that are made with IV as a natural experiment identification strategy could be avoided if researchers used the potential outcomes framework to reason about the IV assumptions. One way to do this is to use the randomized encouragement design as a template for any IV-based natural experiment, as this generally helps the analyst to understand whether IV assumptions are plausible in a given setting. In sum, reasoning about identification assumption is a critical skill.

4 The Design-Based Approach

Throughout the causal inference literature one will invariably notice many references to the importance of design and a general emphasis on the design-based approach. Unfortunately there isn't a widely agreed-upon definition of what it means to use a design-based approach. Dunning (2012) maintains that only natural experiments can be classified as design based.⁵ Imbens (2010, 403) uses a much broader definition, saying that under the design-based approach the analyst places an explicit emphasis on reducing heterogeneity, clarity about identifying assumptions, a concern about endogeneity, and the role of research design.

We might define the design-based approach by saying it is a mode of statistical analysis that emphasizes design rather than statistical modeling. This begs the question of what design is. Rubin (2008, 810) defines design as all contemplating, collecting, organizing, and analyzing of data that takes place prior to seeing any outcome data. Here, I outline a non-exhaustive list of important insights and techniques that have become part of the design-based approach. Each technique, alone or in combination, can be used to bolster the credibility of an identification strategy. These are techniques that allow the analyst to argue that he or she is more likely to distinguish treatment effects from plausible alternatives or biases. As such, these methods can generally be combined with any identification strategy.

4.1 Reducing Heterogeneity

In causal inference, one key challenge is separating possible treatment effects from characteristics of units that may be correlated with treatment status. If the units were exactly identical before treatment, then any differences after treatment could be ascribed to the treatment. The difficulty is that in the social sciences the study units display considerable heterogeneity. Any kind of variability among the study units may be termed heterogeneity. While randomization deals with heterogeneity without eliminating it, there is often reason to reduce heterogeneity in any research design. In randomized experiments the reduction of heterogeneity can occur through blocking before randomization and

⁵Dunning (2012) generally uses the phrase "design-based inference" instead of "design-based approach." I exclusively use the term "design-based approach" to avoid confusion with an older use of the term "design-based inference" used in the literature on survey sampling.

allows for more precise estimation of the treatment effect. In an observational study, reducing heterogeneity often means reducing the sample size to a smaller, more comparable subset.

An example is useful. In a study to understand whether wearing a helmet on a motorcycle reduces the risk of death, Norvell and Cummings (2002) restricted their study to only cases where there were two riders on the motorcycle and one used a helmet but the other rider did not. They reduced heterogeneity by looking at the within-motorcycle pairs instead of simply comparing crashes where one rider had a helmet to other crashes where the rider did not use a helmet. By using the within-pair comparison, they reduced heterogeneity in factors such as road conditions, traffic patterns, and different speeds. Natural experiments often focus on unrepresentative portions of the population where heterogeneity is lower.

One might object to this practice since throwing away data will reduce statistical efficiency. However, efficiency should generally be a secondary concern in observational studies. Why is efficiency a secondary concern in observational studies? The basic insight is from Cochran and Chambers (1965), who demonstrate that if there is a fixed bias that does not decrease as the sample size grows, then as the sample size increases this bias will dominate the mean-squared error for the estimate of the treatment effect. In other words, increasing the sample size can shrink the confidence intervals to a point that excludes the true treatment effect point estimate. In a randomized experiment, where the estimate is known to be unbiased, adding additional observations simply increases power. In an observational study, any additional data that contribute to the heterogeneity may increase bias.

In general the call to reduce heterogeneity arises from differential concerns about sampling uncertainty and uncertainty from unobserved confounding. In observational data, the amount of bias that results from unobserved confounders is a far greater source of uncertainty than uncertainty from a limited sample size. Increasing the sample size, moreover, does nothing to reduce the bias from unobserved confounders. Rosenbaum (2004, 2005a) has analytically demonstrated that reducing unit heterogeneity in observational data reduces sensitivity to bias from unobserved confounders. Reducing unit heterogeneity amounts to restricting the analysis to a more homogeneous subset of the entire data set. One might argue that the concomitant reduction in sample size will reduce the power to detect treatment effects, but this is not the case. Rosenbaum (2004, 2005a) proves that when treatments are nonrandomly assigned, reducing unit heterogeneity reduces *both* sampling variability and sensitivity to bias from unobserved covariates. In short, there are reasons for focusing on small samples where differences across treated and control units are reduced not by statistical means but by the design.

This move to reduce heterogeneity has led to a specific practice in observational studies. Sometimes it is quite difficult to find a control group that we judge to be similar enough to the treated group. In short, the analyst judges that there is too much heterogeneity across the two groups. Often this occurs because there are treated observations that are very different from any of the control units. One solution is to drop the incomparable treated units from the study and restrict the analysis to the subset of the treated units that are comparable. Crump et al. (2009), Rosenbaum (2012), and King, Lucas, and Nielsen (2014) have developed methods for dropping incomparable treated observations. See Zubizarreta et al. (2013) and Keele, Titiunik, and Zubizarreta (2014) for examples of analyses of this type. Importantly, these methods change the estimand. As soon as a single treated unit is dropped, the estimand is some more local version of the ATT. The difficulty is that we no longer have a well-defined estimand. As such, a tension develops between having a well-defined causal estimand and making a credible claim that treated and control groups are comparable in all observable respects. Is this defensible? I would argue that it is.

Identification under the RD design presents a similar dilemma. Strictly speaking, the causal effect is identified exactly at the cutoff, but in practice, we use some subset of observations above and below the cutoff. While there are a number of principled methods for selecting this neighborhood, we are selecting a somewhat arbitrary set of the treated units that are deemed comparable to the controls.⁶ As Rosenbaum (2012) notes, “often the available data do not

⁶See Imbens and Kalyanaraman (2012) and Calonico, Cattaneo, and Titiunik (2013) for recent methods on selecting the neighborhood.

represent a natural population, and so there is no compelling reason to estimate the effect of the treatment on all people recorded in this source of data.” In general, it is not worth holding the estimand inviolate in the face of observable bias. So, researchers have two choices when subjects lack comparability. Give up and declare the identification strategy implausible, or alter the estimand and focus on a subset of the sample where heterogeneity is not a threat. If the analysts adopts the latter strategy, they should be quite clear that the estimand had to change in order to make the identification strategy credible.

4.2 Falsification Tests

Falsification tests come in various forms, but generally focus on testing for treatment effects in places where the analyst knows they should not exist. Causal theories may do more than predict the presence of a causal effect; causal theories may also predict an absence of causal effects. When we find causal effects where they should not be, this is often a sign of hidden confounders and a failure of the identification strategy.

Rosenbaum (2002b) relates a useful example of using a falsification test. In a study of treated and control groups, researchers were interested in whether eating fish contaminated with methylmercury caused chromosomal damage.⁷ In this study, the researchers used a selection on observables identification strategy in forming the treated and control groups, where the treated group was known to have consumed contaminated fish. One way we might understand whether selection on observables is reasonable is to use a falsification test. We cannot prove that selection on observables holds, but we may find clear evidence that it does not hold. In the study, researchers collected data on a number of health-related outcomes, including whether subjects had asthma. There is currently no evidence that methylmercury causes asthma in any form. Researchers could then test for a treatment effect on asthma since it is an outcome known to be unaffected by the treatment. The presence of an effect on asthma would serve as evidence against the selection on observables assumption. That is, a treatment effect on asthma indicates that there is some unobservable difference across the treated and control groups that creates a treatment effect where none should exist. Falsification tests are often used with RD designs. In an RD design, we shouldn’t find that the discontinuity has an effect on any pre-treatment covariates. Falsification tests of this type are often referred to as placebo tests. It is important to emphasize that falsification tests are negative in nature. They provide evidence against the validity of an identification strategy, but no evidence that identification does actually hold.

4.3 Sensitivity Analysis

Sensitivity analyses are another element of a design-based approach. Many sensitivity analyses are based on a partial identification strategy, where bounds are placed on quantities of interest while a key assumption is relaxed. The phrase “sensitivity analysis” is often used informally. Formally a sensitivity analysis is designed to *quantify* the degree to which a key identification assumption must be violated in order for a researcher’s original conclusion to be reversed. A sensitivity analysis provides a quantifiable statement about the plausibility of an identification strategy. If a causal inference is sensitive, a slight violation of the assumption may lead to substantively different conclusions. The first sensitivity analysis explored whether it was possible for an unobserved confounder to explain the leftover variation in lung cancer rates after accounting for the association with smoking (Cornfield et al. 1959). While a sensitivity analysis can be conducted for any identification strategy, most sensitivity analyses focus on the selection on observables assumption (Rosenbaum 1987; Imbens 2003). For many identification strategies, specific forms of sensitivity analysis have not yet been developed.

Briefly, I outline the logic behind one form of sensitivity analysis. Rosenbaum (2002b) has developed a method of bounds to understand whether the selection on observables identification

⁷The original study was conducted by Skerfving et al. (1974).

assumption is sensitive to the presence of a hidden confounder. Under this method, one places bounds on quantities such as the treatment effect point estimate or p -value based on a conjectured level of confounding. That is, the analyst states that he or she thinks the level of the confounding is a given magnitude. For that level of confounding, one can calculate bounds on the treatment effect point estimate. If zero is included in those bounds, a failure of the identification strategy would reverse the study conclusions for that level of confounding. One can vary the level of confounding to observe whether a small or large amount of confounding would reverse the study conclusions.

4.4 *Pattern Specificity*

I conclude this section with one final observation. Statistical results from a single analysis are rarely considered to provide definite proof of a causal relationship. Instead, analysts demonstrate causal relationships by building a multifaceted pattern of evidence. Rosenbaum (2005b) uses the phrase “pattern specificity” to describe the evidence-building process needed in a causal analysis. The concept behind pattern specificity is simple: one should test as many relevant implications of a causal theory as possible. Confirmation of each additional implication strengthens the evidence for a causal effect. Thus, a pattern of specific confirmatory tests provides better evidence than a single test. As Cook and Shadish (1994, 95) write: “Successful prediction of a complex pattern of multivariate results often leaves few plausible alternative explanations.” Under pattern specificity, part of the design is the generation and testing of a large number of hypotheses based on the causal theory. If a series of tests are successful, it lends greater credibility to the causal theory. Many of the techniques described above are often key elements in pattern specificity as one might use falsification tests and sensitivity analysis as part of a single research design.

In this section, I have highlighted the importance of the design-based approach. In general, causal analysis under a design-based approach seeks a plausible identification strategy and then often employs the techniques above to bolster the credibility of that strategy. While none of these techniques in isolation can rule out the presence of hidden bias, they can often increase the credibility of many identification strategies.

5 Tools for Causal Inference

In this final section, I provide an overview of a number of methods that are often used in the analysis of treatment effects. Most of these methods are concerned with estimation of treatment effects and statistical inference for those estimates. That is, once an identification strategy has been selected and the design is complete, the analyst next turns to the estimation of causal effects. A number of new methods have been developed for the estimation of causal effects. I provide little detail on these various methods, as they are covered in much greater depth elsewhere. The appendix contains links and references for software tools available for the methods discussed below.

5.1 *Directed Acyclic Graphs*

One tool that is sometimes applied in the literature on causal inference is that of causal graphs or directed acyclic graphs (DAGs) (Pearl 1995). Unlike the other methods outlined in this section, DAGs are a tool for identification as opposed to statistical analysis. DAGs are often useful for reasoning about causal structure, since they allow us to formalize identification concepts in a graphical manner. From a given graph, we can derive nonparametric identification results and identify which variable or sets of variables are necessary for identification. Pearl (2009a) maintains that DAGs are essential to any causal analysis. A more limited view of DAGs would say that a DAG is meant to represent the analyst’s reasoned view of the causal structure between a set of variables. Once the DAG is written down, it can be defended as a causal representation of a theory. Based on that structure one can then derive whether a causal effect is nonparametrically identified or not. However, in cases where identification conditions are well understood, a DAG may add little to the analysis. That is, in a well-conducted randomized experiment or a good natural

experiment, the design creates such a simple DAG that they are of little use. However, under selection on observables, DAGs can be a useful way to clarify the necessary conditioning set for identification to hold.

5.2 Estimation Methods

The number and variety of statistical methods used in the estimation of causal effects is well beyond the scope of this article. Below, I provide a high-level overview of the methods used. While identification is, strictly speaking, separate from estimation, an emphasis on nonparametric identification tends to influence estimation. When nonparametric identification holds, it implies a valid nonparametric estimator. Thus, if a convincing case can be made for nonparametric identification, in theory nonparametric estimation provides a straightforward way to estimate the identified treatment effect.

What is the problem with straying too far from the implied nonparametric estimator? The danger is that if the analyst selects an overly restrictive method of statistical estimation, estimates of nonparametrically identified causal effects will be biased due to overly restrictive modeling assumptions. For example, assume that selection on observables holds but unit response to treatment is nonlinear. If the analyst applies an estimation method that assumes a linear response to treatment, functional form misspecification may bias the effect such that one might think the treatment is without effect when in fact the effect is simply nonlinear. It would be unfortunate to waste identification due to functional form misspecification. The possibility of bias from functional form misspecification leads to a strong preference for nonparametric or semiparametric estimation methods.⁸ While data or other practical limitations may make nonparametric estimation infeasible, many of the methods used in causal analyses tend to be either nonparametric or semiparametric.

5.2.1 Regression

Here, I use the word “regression” broadly to include not only least squares but also models with nonlinear links such as logistic regression models. The primary use of regression models is to adjust for confounders under selection on observables. However, regression models may be used in conjunction with most of the identification strategies described in this essay. For example, regression-based methods are often used under both the IVs and RD design identification strategies. This illustrates why statistical techniques are secondary to identification strategies. The credibility of the estimator is often a function of the identification strategy, and many methods of estimation have some applicability across different identification strategies.

Many researchers view regression models as estimators of causal effects with suspicion given the strong functional form assumptions needed. Regression models need not be wedded to restrictive functional forms, though. They can be made more flexible through the use of splines or kernel methods (Keele 2008; Hainmueller and Hazlett 2013). Hill, Weiss, and Zhai (2011) and Hill (2011) show how very flexible nonparametric methods that are loosely regression based can be used to estimate causal effects.

Many critiques of regression, however, extend beyond the restrictive functional form. Regression models have been strongly critiqued as a method of the estimation of causal effects (Freedman 2005; Berk 2006). For example, regression models often produce treatment effect estimates based on extrapolation that is not readily observable to the analyst. The basic interpretation of the regression coefficient as a marginal effect can lead to causal interpretations of regression models where identification is questionable. That is, the statement that the β coefficient in a regression model is the amount Y changes for a unit change in X is an implicitly causal statement that is unjustified without careful consideration of the identification strategy.

Regression models, however, also serve auxiliary purposes in a causal analysis. For example, the propensity score is the probability of being exposed to a specific treatment, and they are often used

⁸There are always exceptions. See Angrist and Pischke (2009, chap. 3) for a dissenting view.

in matching or weighting analyses. In both cases, a logistic regression model is typically used to estimate the propensity score and thus is not the estimator of the causal effect, but the regression model serves a key role in the analysis.

5.2.2 Matching

Matching methods are often used in analyses that focus on causal effects. Most frequently, matching is used in conjunction under selection on observables to make treated and control groups identical in terms of observed covariates. Matching is equivalent to a specific form of nonparametric regression. See Angrist and Pischke (2009, 69) for a discussion of the equivalencies. Matching, like regression, has a wide variety of uses across different identification strategies. Often natural experiments based on instruments require statistical adjustment; this form of adjustment can also be done via matching (Rosenbaum 2002a). Recently, matching has been adapted to RD designs (Keele, Titiunik, and Zubizarreta 2014). I credit the more recent popularity of matching to work in economics where matching recovered the estimate from a randomized experiment based on observed covariates (Dehejia and Wahba 1999). This has also led to some confusion, where matching has been mistaken for an identification strategy. See Sekhon (2009) and Arceneaux, Gerber, and Green (2006) for overviews of this confusion. However, it is worth repeating that matching is a statistical technique that is devoid of any identification assumptions. When matching is applied to an IV application, the identification assumptions are completely different from when matching is applied to an application where identifiability is based on selection on observables.

The main attraction of matching is that it is a completely nonparametric form of adjustment. I also think it has advantages in that one can completely customize the form of statistical adjustment. For example, one might dictate very close or exact matches on key variables and looser constraints on covariates that are less important. Balance testing also makes it readily apparent whether matching has succeeded in creating an observably comparable control group for the treated. Matching, however, is simply a tool and cannot compensate for a poor identification strategy. Matching can also be part of the design. For example, matching can be used as a form of blocking in randomized experiments (Greevy et al. 2004). Here, units are made more comparable before treatments are assigned.

5.2.3 Weighting

Besides regression methods and matching, IP weighting is the other major statistical method that has been developed specifically for the estimation of treatment effects (Robins, Rotnitzky, and Zhao 1994; Robins 1999). IP weighting methods can be used to estimate treatment effects in a variety of situations but have seen widespread use in contexts with repeated and time-varying treatments. Glynn and Quinn (2010) provide a useful overview of these methods in a social science context.

Under this method of estimation, the analyst reweights observations to create a pseudo-population where treated and control units are conditionally independent of treatment status. This pseudo-population is created by weighting each unit in the study by the inverse of what is known as the propensity score. I define \mathbf{x} as a matrix of covariates that are thought to be predictive of treatment status, and $e(\mathbf{x}) = P(D_i = 1 | \mathbf{x})$ as the conditional probability of exposure to treatment given observed covariates \mathbf{x} . The quantity $e(\mathbf{x})$ is generally known as the propensity score (Rosenbaum and Rubin 1983). The treatment effect estimate is simply the difference in means across treatment status within the pseudo-population. A number of alternative methods for estimating weights are available, and the estimation of these weights forms an area of active research. IP weighting techniques are also closely identified with what are known as “doubly robust” methods, though double robustness can also be achieved using matching methods (Ho et al. 2007). Double-robust methods model both the treatment assignment mechanism and the outcome. If at least one of these models is correctly specified, the estimate of the ATE will be consistent (Scharfstein, Rotnitzky, and Robins 1999). The double-robust property is no magic bullet since poor estimation of the weights or misspecification of both models may cause bias

(Kang and Schafer 2007). One advantage of IP weighting is that it can also be used to model missingness in the outcomes. Moreover, variance calculations that take into account uncertainty in both the model of treatment and outcome are also straightforward.

5.3 *Inferential Methods*

In the analysis of causal effects, one could easily assume that little changes in terms of statistical tests. For example, in the analysis of an experiment, the usual t -test is typically applied to test whether the ATE is zero. In reality, a subtle change has occurred in the underlying logic of the statistical test. The standard justification for statistical inference is to characterize uncertainty about a random sample from a population. Of course, many experiments are not conducted with representative samples, and yet they can still lead to valid inferences about causal effects for the units under study. Generally in studies of causal effects, the mode of statistical inference is different. Our main source of uncertainty is about whether a causal effect is real or instead a chance outcome due to the stochastic nature of the treatment assignment mechanism. That is, we wish to characterize the probability that an observed treatment effect estimate is large due strictly to chance. The difference in the nature of statistical inference prompted Rubin (1991) to advise analysts to ask: what is your mode of inference?

This question is important since in the study of causal effects, statistical measures of uncertainty depend on how the treatment is assigned. The simplest example arises in randomized experiments. In many randomized experiments, treatments are assigned at the unit level. For example, a GOTV treatment could be assigned to individual-level voters. However, we might instead conduct a group randomized trial, where groups of units are assigned to treatment or control. Under a group RCT, the GOTV treatment might be assigned to households or entire precincts. The difference in assignment mechanisms has implications for measures of statistical uncertainty. If we analyze the group trial as if it were an individual-level trial, the analyst will underestimate statistical uncertainty, since the number of groups is more relevant to calculations of statistical uncertainty than the number of individuals. Thus, it is important to have clarity about how treatments are assigned, since statistical inference directly depends on the treatment assignment mechanism. The mode of inference question becomes more complex outside of experiments since we often do not directly observe how treatments were assigned. In observational data, it is often unclear whether the treatment assignment mechanism operates at a unit or group level, so analysts must carefully consider how to characterize statistical uncertainty. As such, it is important that analysts understand how statistical inference differs when causal effects are the goal.

Statistical inference for treatment effects is typically defined using one of two different frameworks. The first framework is associated with Jerzy Neyman, and the second framework was developed by Ronald Fisher. Here, I briefly point out differences between the two frameworks and discuss why I think it is important to be familiar with both frameworks. Under the Neyman framework, we ask what would be the average outcome if all units were exposed to treatment and how that would compare to the average outcome if all units were exposed to control. The statistical test under the Neyman framework is whether the average causal effect is zero. In the Fisherian framework, we test what is known as the sharp null hypothesis. Under the sharp null hypothesis, the analyst tests whether the treatment effect is zero for every unit. In potential outcomes notation, if the sharp null hypothesis holds, then $Y_{1i} = Y_{0i}$ for every i . In the Fisherian approach, there is no way to test the null hypothesis that the average effect is zero (Imbens and Rubin 2015).⁹ This might strike some readers as a major drawback, since this would seem to be a very restrictive null hypothesis. One advantage of testing average effects is that we can accommodate heterogeneous responses to treatment. That is, under a test of the average effect, the units can have some mix of positive and negative responses to treatment.

⁹There are a number of other features that are unique to the Fisherian framework, including that it can be used as a method of estimation. Keele, McConaughy, and White (2012) provide a basic overview of Fisher's approach.

However, only testing for average causal effects has pathologies of its own. Take an example from Imbens and Rubin (2015). Let's say that $Y_{0i} = 2$. For one-third of the units in the study, the treatment effect is 2, but for two-thirds of the units, the effect is -1 . Here, the average effect is zero, but the sharp null is not. Again the mode of inference matters, in that we might detect an effect with one mode of inference but miss it with another. One particular strength of the Fisherian framework is that it can accommodate a wide variety of tests about quantities other than averages. Thus far, I have described estimands only as averages. However, there is nothing that specifically implies that a treatment will only change the middle of the treated and control distributions as summarized by the average. In the most extreme example, the treatment might only change the variance of the treated distribution. Under the Fisherian framework, we can apply the Kolmogorov–Smirnov statistic, which tests the maximum discrepancy in the empirical cumulative distribution functions (CDFs) and can detect differences in any of the moments of the distribution. The Fisherian framework has also been extended in other fruitful ways. It serves as the basis for one common method of sensitivity analysis (Rosenbaum 2002b). Bowers, Fredrickson, and Panagopoulos (2013) use it to analyze empirical applications with treatment spillovers.

I would argue that analysts need to be familiar with both frameworks. A clear understanding of both is useful in two ways. One is that it clarifies how the mode of statistical inference matters. Under the Fisherian framework, it is obvious how the mode of inference changes depending on the assignment mechanism. Moreover, it allows for testing quantities other than average effects. The Neyman framework, however, allows for tests of average effects, which are at a minimum a useful starting point. This framework also accommodates sampling from populations, which arises when randomized experiments are conducted with random samples from populations. Generally, after an examination of average effects, analysts should consider whether other features of the treated and control distributions differ and test for such differences.

6 Discussion

The reader may notice that this essay is heavily tilted toward identification rather than on the intricacies of matching methods or the relative merits of doubly robust estimators. It is not because estimation and inference aren't important, but it is due to the fact that no statistical method can save a poor identification strategy. Many of the pathologies in the statistical analysis of causal effects stem from confusion over the separate roles of identification and estimation. Understanding this distinction provides an important check on what it is that analysts think they can learn from data. Much of the language in statistics has long obscured the importance of assumptions. To say that a model is unbiased when correctly specified is a true statement, and yet seriously understates how difficult it can be to achieve the correct specification when the goal is estimation of causal effects. An understanding of what it means for something to be correctly specified (i.e., identified) reveals the limits of what can generally be learned from data about causal effects, especially with observational data. Moreover, it reveals that complex statistical estimators may do nothing to aid an identification strategy. Causal inference, particularly in relation to topics where randomized experiments are impossible, will probably remain a difficult task that requires a series of different identification strategies across a number of different contexts before conclusions can be reached.

To that end, one goal in this essay was to illuminate an important paradox that lies within the statistics of causal inference. This paradox occurs in the fact that the most credible causal inferences require the least amount of statistical analysis. In fact, when a causal inference is credible, most of the work will have been done before the outcome data are collected. If the analyst has taken the time to develop a randomized treatment assignment mechanism and reduce or eliminate noncompliance and attrition, often the analysis of the causal effect is reduced to a simple contrast in measures of distributional location. The analyst is successful at identifying the causal effect not because of the complex statistical methods that are applied to the data, but due to the effort in developing a design before data are collected. A quote by Fisher (1938) is instructive on this point: "To consult the statistician after an experiment is finished is often merely to ask him to conduct a post-mortem examination. He can perhaps say what the experiment died of."

To conclude, I reflect briefly on what remains to be done. I think one open question is the extent to which everyday statistical practice will absorb the view of causality presented in this essay. Much of this view stands at odds with what might be called standard statistical practice. It is certainly different from what I was taught in graduate school. A causal inference approach based on identification tends to take a rather skeptical view of statistical analyses based on selection on observables. Many methodologists who focus on causal inference also take a dim view of regression models, which remain by far and away the most commonly used statistical method. I think there is much work to be done simply in terms of communication to the larger discipline. This essay is meant to serve as one initial step in that direction. I think the identification approach can lead to much greater credibility in social science research, but only if applied analysts understand the value of this framework.

There are also many avenues of active research. Many applications in political science could be considered to have dynamic treatments that repeat over time. The literature on dynamic treatments in biostatistics has developed mostly in response to a specific type of clinical trial that bears little resemblance to social science applications. There is much that could be done in terms of developing dynamic methods for social science contexts. Causal inference in the presence of interference across units will continue to be an important avenue for future research given that social interaction will invariably be a feature of many political processes. Finally, I also think that, given the difficulty of conducting randomized trials in many areas of political science, the partial identification strategy deserves greater use. Current partial identification strategies tend to be quite general and very conservative. There is much work to be done on the development of partial identification strategies that can help us understand whether inferences outside of experiments are credible. Partial identification lacks the certainty of point estimates, but allows one to clearly communicate how estimates change as assumptions are relaxed. In conclusion, our understanding of the role that statistics play in causal inference has changed greatly. Causal inference is difficult, but progress can be made.

Conflict of interest statement. None declared.

References

- Abadie, Alberto, Alexis Diamond, and Jens Hainmueller. 2010. Synthetic control methods for comparative case studies: Estimating the effect of California's tobacco control program. *Journal of the American Statistical Association* 105(490):493–505.
- Abadie, Alberto, and Javier Gardeazabal. 2003. The economic costs of conflict: A case study of the Basque country. *American Economic Review* 93(1):112–32.
- Angrist, Joshua D., Guido W. Imbens, and Donald B. Rubin. 1996. Identification of causal effects using instrumental variables. *Journal of the American Statistical Association* 91(434):444–55.
- Angrist, Joshua D., and Jörn-Steffen Pischke. 2009. *Mostly harmless econometrics*. Princeton, NJ: Princeton University Press.
- . 2010. The credibility revolution in empirical economics: How better research design is taking the con out of econometrics. *Journal of Economic Perspectives* 24(2):3–30.
- Arceneaux, Kevin, Alan S. Gerber, and Donald P. Green. 2006. Comparing experimental and matching methods using a large-scale voter mobilization study. *Political Analysis* 14(1):37–62.
- Balke, Alexander, and Judea Pearl. 1997. Bounds on treatment effects from studies with imperfect compliance. *Journal of the American Statistical Association* 92(439):1171–1176.
- Barnow, B. S., G. G. Cain, and A. S. Goldberger. 1980. Issues in the analysis of selectivity bias. In *Evaluation studies*, eds. E. Stromsdorfer and G. Farkas, Vol. 5, 43–59. San Francisco: Sage Publications.
- Berk, Richard A. 2006. *Regression analysis: A constructive critique*. Thousand Oaks, CA: Sage Publications.
- Bound, J., D. A. Jaeger, and R. M. Baker. 1995. Problems with instrumental variables estimation when the correlation between the instruments and the endogenous explanatory variable is weak. *Journal of the American Statistical Association* 90(430):443–50.
- Bowers, Jake, Mark M. Fredrickson, and Costas Panagopoulos. 2013. Reasoning about interference between units: A general framework. *Political Analysis* 21(1):97–124.
- Calonico, Sebastian, Matias Cattaneo, and Rocio Titiunik. 2013. Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica* 82(6):2295–326.
- Campbell, Donald T., and Julian C. Stanley. 1963. *Experimental and quasi-experimental designs for research*. Chicago: Rand McNally.

- Cattaneo, Matias, Brigham Frandsen, and Rocio Titiunik. 2014. Randomization inference in the regression-discontinuity design: An application to party advantages in the U.S. Senate. *Journal of Causal Inference*. Unpublished manuscript.
- Caughey, Devin, and Jasjeet S. Sekhon. 2011. Elections and the regression discontinuity design: Lessons from close U.S. House races, 1942–2008. *Political Analysis* 19(4):385–408.
- Cochran, William G., and S. Paul Chambers. 1965. The planning of observational studies of human populations. *Journal of Royal Statistical Society, Series A* 128(2):234–65.
- Cook, T. D., and W. R. Shadish. 1994. Social experiments: Some developments over the past fifteen years. *Annual Review of Psychology* 45:545–80.
- Cook, Thomas D., William R. Shadish, and Vivian C. Wong. 2008. Three conditions under which experiments and observational studies produce comparable causal estimates: New findings from within-study comparisons. *Journal of Policy Analysis and Management* 27(4):724–50.
- Cornfield, J., W. Haenszel, E. Hammond, A. Lilienfeld, M. Shimkin, and E. Wynder. 1959. Smoking and lung cancer: Recent evidence and a discussion of some questions. *Journal of National Cancer Institute* 22:173–203.
- Crump, Richard K., V. Joseph Hotz, Guido W. Imbens, and Oscar A. Mitnik. 2009. Dealing with limited overlap in estimation of average treatment effects. *Biometrika* 96(1):187–99.
- Dawid, A. Philip. 2000. Causal inference without counterfactuals. *Journal of the American Statistical Association* 95(450):407–24.
- Dehejia, Rajeev, and Sadek Wahba. 1999. Causal effects in non-experimental studies: Re-evaluating the evaluation of training programs. *Journal of the American Statistical Association* 94(448):1053–1062.
- Ding, Peng, and Luke W. Miratrix. 2015. To adjust or not to adjust? Sensitivity analysis of M-bias and butterfly-bias. *Journal of Causal Inference* 3(1):41–57.
- Dunning, Thad. 2012. *Natural experiments in the social sciences: A design-based approach*. Cambridge, UK: Cambridge University Press.
- Elwert, Felix, and Christopher Winship. 2014. Endogenous selection bias: The problem of conditioning on a collider variable. *Annual Review of Sociology* 40(1):31–53.
- Fisher, R. A. 1938. Presidential address. *Sankhya: The Indian Journal of Statistics* 4(1):14–7.
- Freedman, D. A. 2005. Linear statistical models for causation: A critical review. *Encyclopedia of Statistics in Behavioral Science*.
- Gerber, Alan S., and Donald P. Green. 2012. *Field experiments: Design, analysis, and interpretation*. New York: Norton.
- Glynn, Adam N., and Kevin M. Quinn. 2010. An introduction to the augmented inverse propensity weighted estimator. *Political Analysis* 18(1):36–56.
- Gordon, Sanford C. 2011. Politicizing agency spending authority: Lessons from a bush-era scandal. *American Political Science Review* 105(4):717–34.
- Greevy, Robert, Bo Lu, Jeffery H. Silber, and Paul Rosenbaum. 2004. Optimal multivariate matching before randomization. *Biostatistics* 5(2):263–75.
- Hahn, Jinyong, Petra Todd, and Wilbert van der Klaauw. 2001. Identification and estimation of treatments effects with a regression-discontinuity design. *Econometrica* 69(1):201–9.
- Hainmueller, Jens, and Chad Hazlett. 2014. Kernel regularized least squares: Reducing misspecification bias with a flexible and interpretable machine learning approach. *Political Analysis* 22(2):143–168.
- Hansford, Thomas G., and Brad T. Gomez. 2010. Estimating the electoral effects of voter turnout. *American Political Science Review* 104(2):268–88.
- Hernán, Miguel A., and Tyler J. VanderWeele. 2011. Compound treatments and transportability of causal inference. *Epidemiology* 22(3):368–77.
- Hidalgo, Daniel F., and Jasjeet S. Sekhon. 2011. Causation. In *International Encyclopedia of Political Science*, eds. Bertrand Badie, Dirk Berg-Schlosser, and Leonardo Morlino, 203–10. Thousand Oaks, CA: Sage Publications.
- Hill, Jennifer, Christopher Weiss, and Fuhua Zhai. 2011. Challenges with propensity score strategies in a high-dimensional setting and a potential alternative. *Multivariate Behavioral Research* 46(3):477–513.
- Hill, Jennifer L. 2011. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics* 20(1):217–40.
- Ho, Daniel E., Kosuke Imai, Gary King, and Elizabeth A. Stuart. 2007. Matching as nonparametric preprocessing for reducing model dependence in parametric causal inference. *Political Analysis* 15(3):199–236.
- Holland, Paul W. 1986. Statistics and causal inference. *Journal of the American Statistical Association* 81(396):945–60.
- . 1988. Causal inference, path analysis, and recursive structural equation models. *Sociological Methodology* 18:449–84.
- Imai, Kosuke, Luke Keele, Dustin Tingley, and Teppei Yamamoto. 2011. Unpacking the black box of causality: Learning about causal mechanisms from experimental and observational studies. *American Political Science Review* 105(4):765–89.
- Imbens, Guido W. 2003. Sensitivity to exogeneity assumptions in program evaluation. *American Economic Review Papers and Proceedings* 93(2):126–32.
- . 2010. Better LATE than nothing: Some comments on Deaton (2009) and Heckman and Urzua (2009). *Journal of Economic Literature* 48(2):399–423.
- Imbens, Guido W., and Donald B. Rubin. 2015. *Causal inference for statistics, social, and biomedical sciences: An introduction*. Cambridge, UK: Cambridge University Press.

- Imbens, Guido W., and Karthik Kalyanaraman. 2012. Optimal bandwidth choice for the regression discontinuity estimator. *Review of Economic Studies* 79(3):933–59.
- Kang, Joseph D. Y., and Joseph L. Schafer. 2007. Demystifying double robustness: A comparison of alternative strategies for estimating a population mean from incomplete data. *Statistical Science* 22(4):523–39.
- Keele, Luke. 2008. *Semiparametric regression for the social sciences*. Chichester, UK: Wiley and Sons.
- Keele, Luke J., Corrine McCaughy, and Ismail K. White. 2012. Strengthening the experimenter's toolbox: Statistical estimation of internal validity. *American Journal of Political Science* 56(2):484–99.
- Keele, Luke J., and William Minozzi. 2012. How much is Minnesota like Wisconsin? Assumptions and counterfactuals in causal inference with observational data. *Political Analysis* 21(2):193–216.
- Keele, Luke, Rocío Titiunik, and José Zubizarreta. 2014. Enhancing a geographic regression discontinuity design through matching to estimate the effect of ballot initiatives on voter turnout. *Journal of the Royal Statistical Society, Series A* 178(1):223–39.
- King, Gary, Christopher Lucas, and Richard Nielsen. 2014. The balance-sample size frontier in matching methods for causal inference. Unpublished Manuscript.
- Lee, David S. 2008. Randomized experiments from non-random selection in U.S. House elections. *Journal of Econometrics* 142(2):675–97.
- . 2009. Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *Review of Economic Studies* 76(3):1071–102.
- Lee, David S., and Thomas Lemieux. 2010. Regression discontinuity designs in economics. *Journal of Economic Literature* 48(2):281–355.
- Lyall, Jason. 2009. Does indiscriminate violence incite insurgent attacks? Evidence from Chechnya. *Journal of Conflict Resolution* 53(3):331–62.
- Manski, Charles F. 1990. Nonparametric bounds on treatment effects. *American Economic Review Papers and Proceedings* 80(2):319–23.
- . 1995. *Identification problems in the social sciences*. Cambridge, MA: Harvard University Press.
- . 2007. *Identification for prediction and decision*. Cambridge, MA: Harvard University Press.
- Matzkin, Rosa L. 2007. Nonparametric identification. *Handbook of Econometrics* 6:5307–68.
- Mebane, Walter R., and Paul Poast. 2013. Causal inference without ignorability: Identification with nonrandom assignment and missing treatment data. *Political Analysis* 22(2):169–82.
- Morgan, Stephen L., and Christopher Winship. 2014. *Counterfactuals and causal inference: Methods and principles for social research*. 2nd ed. New York: Cambridge University Press.
- Norvell, Daniel C., and Peter Cummings. 2002. Association of helmet use with death in motorcycle crashes. *American Journal of Epidemiology* 156(5):483–87.
- Pearl, Judea. 1995. Causal diagrams for empirical research. *Biometrika* 82(4):669–710.
- Pearl, Judea. 2001. Direct and indirect effects. Proceedings of the Seventeenth Conference on Uncertainty in Artificial Intelligence. San Francisco, CA: Morgan Kaufmann Publishers.
- Pearl, Judea. 2009a. *Causality: models, reasoning, and inference*. 2nd ed. New York: Cambridge University Press.
- . 2009b. Letter to the editor. *Statistics in Medicine* 28:1415–1416.
- . 2010. On the consistency rule in causal inference: Axiom, definition, assumption, or theorem? *Epidemiology* 21(6):872–5.
- Poe, Steven C., and C. Neal Tate. 1994. Repression of human rights to personal integrity in the (1980s): A global analysis. *American Political Science Review* 88(04):853–72.
- Robins, James M. 1997. Causal inference from complex longitudinal data. *Latent variable modeling and applications to causality*, 69–117. New York: Springer.
- . 1999. Marginal structural models versus structural nested models as tools for causal inference. In *Statistical methods in epidemiology: The environment and clinical trials*, eds. E. Halloran and D. Berry, 951–34. New York: Springer-Verlag.
- Robins, James M., Andrea Rotnitzky, and Lue Ping Zhao. 1994. Estimation of regression coefficients when some regressors are not always observed. *Journal of the American Statistical Association* 89(427):846–66.
- Robins, J. M. 2003. Semantics of causal DAG models and the identification of direct and indirect effects. In *Highly structured stochastic systems*, eds. P. J. Green, N. L. Hjort, and S. Richardson, 70–81. Oxford: Oxford University Press.
- Rosenbaum, Paul R. 1984. The consequences of adjusting for a concomitant variable that has been affected by the treatment. *Journal of the Royal Statistical Society, Series A* 147(5):656–66.
- . 1987. Sensitivity analysis for certain permutation inferences in matched observational studies. *Biometrika* 74(1):13–26.
- . 2002a. Covariance adjustment in randomized experiments and observational studies. *Statistical Science* 17(3):286–387.
- . 2002b. *Observational studies*. 2nd ed. New York: Springer.
- . 2004. Design sensitivity in observational studies. *Biometrika* 91(1):153–64.
- . 2005a. Heterogeneity and causality: Unit heterogeneity and design sensitivity in observational studies. *American Statistician* 59(2):147–52.
- . 2005b. Observational study. In *Encyclopedia of statistics in behavioral science*, eds. Brian S. Everitt and David C. Howell, Vol. 3, 1451–1462. Chichester, UK: John Wiley and Sons.
- . 2010. *Design of observational studies*. New York: Springer-Verlag.

- . 2012. Optimal matching of an optimally chosen subset in observational studies. *Journal of Computational and Graphical Statistics* 21(1):57–71.
- Rosenbaum, Paul R., and Donald B. Rubin. 1983. The central role of propensity scores in observational studies for causal effects. *Biometrika* 76(1):41–55.
- Rubin, Donald B. 1974. Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology* 6:688–701.
- . 1986. Which ifs have causal answers. *Journal of the American Statistical Association* 81(396):961–62.
- . 1991. Practical implications of modes of statistical inference for causal effects and the critical role of the assignment mechanism. *Biometrics* 47(4):1213–34.
- . 2008. For objective causal inference, design trumps analysis. *Annals of Applied Statistics* 2(3):808–40.
- Scharfstein, Daniel O., Andrea Rotnitzky, and James M. Robins. 1999. Adjusting for nonignorable drop-out using semiparametric nonresponse models. *Journal of the American Statistical Association* 94(448):1096–1120.
- Sekhon, Jasjeet S. 2009. Opiates for the matches: Matching methods for causal inference. *Annual Review of Political Science* 12:487–508.
- Sekhon, Jasjeet S., and Rocío Titiunik. 2012. When natural experiments are neither natural nor experiments. *American Political Science Review* 106(1):35–57.
- Sinclair, Betsy, Margaret McConnell, and Donald P. Green. 2012. Detecting spillover in social networks: Design and analysis of multilevel experiments. *American Journal of Political Science* 56(4):1055–1069.
- Skerfving, S., K. Hansson, C. Mangs, J. Lindsten, and N. Ryman. 1974. Methylmercury-induced chromosome damage in man. *Environmental Research* 7(1):83–98.
- Sovey, J. Allison, and Donald P. Green. 2011. Instrumental variables estimation in political science: A readers' guide. *American Journal of Political Science* 55(1):188–200.
- Tchetgen, Eric J. Tchetgen, and Tyler J. VanderWeele. 2012. On causal inference in the presence of interference. *Statistical Methods in Medical Research* 21(1):55–75.
- van der Laan, Mark J., Thaddeus J. Haight, and Ira B. Tager. 2005. "van der Laan et al. respond to 'Hypothetical interventions to define causal effects'". *American Journal of Epidemiology* 162(7):621–22.
- Zubizarreta, José R., Dylan S. Small, Neera K. Goyal, Scott Lorch, and Paul R. Rosenbaum. 2013. Stronger instruments via integer programming in an observational study of late preterm birth outcomes. *Annals of Applied Statistics* 7(1):25–50.